CMP-6048A Advanced Programming

Project Report - 13 January 2025

Vec3 Maths Bytecode Interpreter

Group members:
Jamie Wales, Jacob Edwards.

School of Computing Sciences, University of East Anglia

Version 2.0

Abstract

Vec3 is a bytecode interpreted maths language complete with a GUI, plotting, full static type inference and various maths functions.

The language is designed to be simple to use and understand, with a focus on strict mathematical expressions and plotting, but with more powerful constructs such as recursive bindings, first class functions and static type inference.

The language and GUI are written in F#, using Avalonia(**?**) for the GUI and ScottPlot(**?**) for plotting.

Chapter 1

Introduction

The introduction should be brief and comprise the following:

1.1 Project statement

1.2 Aims and objectives

Chapter 2

Background

Chapter 3

Development History

3.1 Sprint 1: Basic expressions

This sprint focused on implementing a lexer and parser for the language, with precedence rules for the arithmetic operators, parsed with Pratt parsing.

3.1.1 Grammar in BNF

```
<expr> ::= <term> | <term> "+" <expr> | <term> "-" <expr>
<term> ::= <factor> | <factor> "*" <term> | <factor> "/" <term>
<factor> ::= <number> | "(" <expr> ")"
<number> ::= <int> | <float>
<int> ::= <digit> | <digit> <int>
<float> ::= <int> "." <int>
<digit> ::= "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
```

3.2 Sprint 2: Variable assignment

In this sprint we added variable assignment to the language, with the ability to bind an expression to a variable name.

3.2.1 Grammar in BNF

```
<stmtlist> ::= <stmt> | <stmt> <stmtlist>
<stmt> ::= <expr> | "let" <identifier> "=" <expr>

<expr> ::= <term> | <term> "+" <expr> | <term> "-" <expr>
<term> ::= <factor> | <factor> "*" <term> | <factor> "/" <term> | <factor> "%" <term>
<factor> ::= <number> | <identifier> | "(" <expr> ")" | <factor> "^" <factor>

<number> ::= <int> | <float>

<int> ::= <digit> | <digit> <int>
<float> ::= <int> "." <int>
<digit> ::= "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"

<identifier> ::= <letter> | <letter> <identifier>
<letter> ::= "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h" | "i" | "j" | "k" | "l" | "m"
```

3.3 Sprint 3: Interpreter

In this sprint a basic interpreter was implemented, with the ability to evaluate expressions and variable bindings.

We used a simple environment to store variable bindings, and a recursive evaluation function to evaluate expressions.

It was a REPL style interpreter, where the last expression of a statement list was evaluated and printed.

3.3.1 Grammar in BNF

3.4 Sprint 4: Functions

In this sprint we added the ability to define and call functions, with a simple lambda syntax of the form (*args*) -> *expr* and function calls of the form *funcName / lambda(args)*.

Call by value semantics were used, with a new environment created for each function call.

3.4.1 Grammar in BNF

```
<stmtlist> ::= <stmt> | <stmt> <stmtlist>
<stmt> ::= <expr> | "let" <identifier> "=" <expr>
<expr> ::= <term> | <term> "+" <expr> | <term> "-" <expr>
<term> ::= <factor> | <factor> "*" <term> | <factor> "/" <term> | <factor> "%" <term>
<factor> ::= <number> | <identifier> | "(" <expr> ")" | <factor> "^" <factor> | <identifi
<lambda> ::= "(" <exprlist> ")" "->" <expr>

<number> ::= <int> | <float>
<int> ::= <digit> | <digit> <int>
<float> ::= <int> "." <int>

<digit> ::= "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
<identifier> ::= <letter> | <letter> <identifier>
<letter> ::= "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h" | "i" | "j" | "k" | "l" | "m"
<exprlist> ::= <expr> | <expr> "," <exprlist>
```

3.5 Sprint 5: Static type checking

In this sprint we added static type checking to the language, with a simple type inference system based on Hindley-Milner.

The concept of types was introduced, with the types `Int`, `Float`, `Bool`, `Function` and `Never`.

3.5.1 Grammar in BNF

3.6 Sprint 6: Bytecode

In this sprint the interpreter was rewritten to use a bytecode interpreter, with a stack based virtual machine as well as a simple bytecode compiler, allowing for more efficient evaluation of expressions.

3.6.1 Grammar in BNF

3.7 Sprint 7: GUI

A simple GUI was developed in order to allow easier testing of the language, with a text box for input and output and a decompiler output for debugging.

3.7.1 Grammar in BNF

3.8 Sprint 8: Plotting

In this sprint we added the ability to plot lists of points, with a simple plotting function that took a list of x coordinates and a list of y coordinates and plotted them on a graph using **?**.

3.8.1 Grammar in BNF

3.9 Sprint 9: Maths Functions

In this sprint we added a number of maths functions to the language, including `sin`, `cos`, `tan`, `asin`, `acos` and others, and added the ability to plot functions, both built in and user defined.

3.9.1 Grammar in BNF

3.10 Sprint 9: Optimisation

In this sprint we added a simple optimisation pass to the bytecode compiler, which removed unnecessary stack operations and combined constant expressions.

3.10.1 Grammar in BNF

Chapter 4

Final deliverable

4.1 Final BNF

4.2 Final GUI

4.3 Notable Features

4.4 Lexer

Initial lexer design was based on a simple regular expression based lexer, but this was later replaced with a more functional approach using pattern matching on the input string.

The reason for this change was that the regular expression based lexer was difficult to extend and maintain due to the lack of type safety. For example if we had a more general regex called before a more specific one, the more general one would always match first, even if the more specific one should have matched.

This was solved by using a more functional approach, where the type system of F# would inform us if a case would never be matched due to the order of the cases or otherwise, preventing a class of easily overlooked errors during development.

The lexer is now implemented as a recursive pattern matching function that takes a string and returns a list of tokens, complete with their lexeme and position in the input string.

Lexer errors are also accumulated in a list of type *LexerError*, which are displayed to the user in the GUI.

Something of note is that the lexer parses numbers itself, rather than passing them to the parser as strings.

Additionally, due to the permittance of user defined operators, the lexer makes special considerations when lexing special characters, as the distinction between a built-in operator (with precedence) and a user defined operator (currently without precedence) is made during lexing.

Furthermore, both block comments (/* */) and line comments (//) are handled by the lexer by ignoring the contents of the comment. In future, it may be interesting represent comments as a token in the AST, allowing for systems such as documentation generation or formatting to be implemented.

4.5 Parser

The parser is implemented using Pratt parsing(**?**), which is a top-down operator precedence parsing method that allows for easy extension and modification of the grammar.

It works by assigning a precedence to each token, as well as functions specifying how to parse the token when encountering it in a prefix, infix or postfix position.

For example, take the expression $2 + 3 * 4$.

The parser would first encounter the number *2*, which has a precedence of 0 and a prefix function that simply returns the number.

Thus, the current state of the parser is 2.

The parser would then encounter the operator *+*, which has a precedence of 1 and a left associative infix function that takes the left hand side and the right hand side and returns a binary expression node.

The parser then attempts to parse the right hand side of the operator with a precedence level higher than the plus operator, as Pratt parsing must ensure that higher precedence operations (such as multiplication) are parsed first.

The parser would then encounter the number *3*, which again is treated as a literal and returned.

The parser then encounters the operator ***, which has a precedence of 2 (higher than the plus operator) and as such the parser cannot yet resolve the *+* operator; it must handle the higher precedence multiplication operator first.

The parser saves the left hand side (the number 3) and then parses the right hand side of the multiplication operator using a precedence level higher than the multiplication operator.

It encounters the number *4*, which is returned as a literal.

The parser then returns the binary expression node for the multiplication operator, with the left hand side being the number 3 and the right hand side being the number 4.

The parser then returns to the plus operator, which can now be resolved as the left hand side is the number 2 and the right hand side is the result of the multiplication operator.

This is a simple example, but Pratt parsing can handle more complex expressions with ease, such as nested expressions and function calls.

Using Pratt parsing has improved the extensibility of the parser, as adding new operators or changing the grammar is as simple as adding a new case to the parser.

A slight limitation is during ambiguity, such as the *(* symbol, which can be used for a grouping, a lambda definition, a tuple or a unit type when encountered in the prefix position.

This is resolved by using a state machine approach, where the parser can move around the state at will, allowing lookahead and backtracking in order to reach a point where the ambiguity is resolved.

In order to simplify the Virtual Machine4.10, the parser parses all binary and unary operations as function calls, with the operator as the function name.

In order to make type inference simpler for operators that are overloaded for both unary and binary operations (such as the - operator), the operator keeps track of the manner in which it is called (unary or binary) and returns the appropriate AST node, allowing for easier type inference (as the names of the otherwise overloaded functions are different), and simplifies the bytecode generation process.

This idea could possibly be extended to allow other overloaded function names (with varying numbers of arguments).

4.6 AST

The AST of the language is represented as a list of statements, where a statement is either expression, a variable assignment or an other statement type.

The AST is typed (after type inference4.7) in order to allow for easier optimisation and bytecode generation.

4.7 Type Inference

Vec3 is a statically typed language, with full type inference. The type inference algorithm is based on Hindley-Milner type inference(Sulzmann, 2000), with some modification to

support the non-ML style syntax, and extended to support *row polymorphism*(Morris and McKinna, 2019) (4.7.3), *gradual typing*(Garcia et al., 2016) (4.7.2), *recursive bindings* (4.7.4), *vector length encoding* (4.7.6) and a seemingly unique method of supporting ad-hoc polymorphism named *constraints* (4.7.5).

The reason for implementing strong type inference due to the *Semantic Soundness Theory*(Timany et al., 2024), which states that a *well-typed program cannot go wrong*.

Of course this is not strictly true in practice due to external factors, but it is certainly true that strong typing rules out a large class of errors, most of which human, and as such it is a valuable tool for a maths language to have as the user is less likely to make trivial mistakes.

Another thing to note is that in order to make the language more intuitive to use, the **integer** type will coerce into any other number type (**float, rational, complex**).

This allows for expressions such as 5.0^5 type checking successfully, with the result being a float.

4.7.1 Type Inference Algorithm

The algorithm used to infer types is based on Algorithm W(Milner, 1978). The general idea is to assign the widest type possible for a given node in the AST, which is generally a *type variable*, which is a type used to represent a type that can be unified with any other type (a generic type). The node's children are then inferred, and the types of the children are unified with the parent node. If the types cannot be unified, then the program is ill-typed. Unification is the process of finding the most general type that can be assigned to two types, and is a key part of the algorithm.

For example, unifying *int* and *int* would result in *int*, as this is the most general type that can be assigned to both.

Contrasting this, unifying int and a type variable *a*, would result in *int*, and then the type variable *a* would have to be substituted with *int* throughout the program (because *int* is the most general type that can be assigned to *a*).

It works bottom-up as only a few types are known at the start, such as the types of literals and the types of built-in functions.

Algorithm Implementation   A simplified version of the algorithm, with some details omitted for brevity, is shown in Algorithm 1.

As shown, it is an incredibly simple yet powerful algorithm, and is the basis for many modern type inference algorithms, such as that of F# and OCaml.

Bindings   Generally in implementations of *Algorithm W*, after type inference for a given binding has taken place a process known as *generalisation* occurs. This is the process of replacing all type variables in the type of the binding with *forall* quantifiers, which is a way of saying that the type is polymorphic and can therefore be instantiated with any type.

However, this was not necessary in our implementation as we don't specialise bindings during the instantiation of types (such as during calls), we simply infer the type of the call and check it against the type of the binding, so generalisation is not necessary.

4.7.2 Gradual Typing

Gradual typing is a type system that allows for the gradual transition from dynamic typing to static typing. This is useful in a language like Vec3 as it allows for the user to write code without having to worry about types allowing for quick prototyping, but then add types later to ensure correctness.

Users have the option of adding types to their code in the form $let\ x : int = 5$, and the type inference algorithm will check that the type of the expression matches the type given.

Algorithm 1 Type Inference Algorithm

```
 1: function unify(type1, type2)
 2:     if type1 is type variable then
 3:         type1 ← type2
 4:     end if
 5:     if type2 is type variable then
 6:         type2 ← type1
 7:     end if
 8:     if type1 is function type and type2 is function type then
 9:         unify paramTypes
10:         unify returnTypes
11:     end if
12:     if type1 is not equal to type2 then
13:         error
14:     end if
15: end function
16: function infer(expr, env)
17:     if expr is literal then
18:         return type of literal
19:     end if
20:     if expr is variable then
21:         T ← lookup variable env
22:         return type
23:     end if
24:     if expr is function call then
25:         funcType ← infer function
26:         argTypes ← infer arguments
27:         funcType ← unify paramTypes argTypes
28:         returnType ← return type of funcType
29:         return returnType
30:     end if
31:     if expr is binding then
32:         bodyType ← infer body
33:         env ← add binding bodyType environment
34:         return bodyType
35:     end if
36:     if expr is lambda then
37:         argTypes ← new type variables
38:         bodyEnv ← add arguments to environment
39:         bodyType ← infer body with bodyEnv
40:         funcType ← argTypes + bodyType
41:         return funcType
42:     end if
43: end function
```

The type *any* can also be used, which represents a dynamic type that can be unified with any other type. This disables the safety guarantees of the type system, but can be useful as mentioned above for quick prototyping.

One thing to note however is that the *any* type is infectious, meaning that if a type is inferred to be *any* then the type of the parent node will also be *any*.

4.7.3 Row Polymorphism

Row polymorphism is a form of polymorphism that allows for the definition of functions that operate on records with a certain set of fields, but can also operate on records with additional fields.

It can be considered both a form of structural typing like that of TypeScript(Bierman et al., 2014), and a form of subtyping.

For example, consider the following function:

```
let f = (x) -> x.a
```

This function takes a record with a field *a* and returns the value of that field.

Now consider the following record:

```
let r = {a = 5, b: int = 6}
```

The function *f* can be called with *r* as an argument as *r* has a field *a*, and the function will return 5.

This is a powerful feature as it allows for the definition of functions that operate on a wide range of records.

The reference algorithm given by Morris and McKinna (2019) was used as a basis for the implementation of row polymorphism in Vec3, however without record restriction as it was not necessary for the language.

The algorithm works by assigning a *row variable* to each record type, which is a type variable that represents the fields of the record, where a record is represented in the type system as an extension of another record, or the empty record.

The algorithm then unifies the row variables of the record types, and if the unification is successful then the function can be called with the record.

A reference implementation of the algorithm is given in Algorithm 2.

Algorithm 2 Row Polymorphism Algorithm

With some creativity, row polymorphism can be used to represent semi-algebraic data types or tagged unions. For example, consider the built-in *on* function (used to add event listeners for shapes)4.12:

```
on(shape, Keys.Down, (state) -> ...)
```

The function expects a shape reference, a key, and a function that takes a state.

The implementation of the keys record is hidden from the user, but could well be implemented as a record with a field for each key, where each key is a record that contains a field *43hr4h54j3* (a unique identifer for the keys record) and the *on* function could have a type of:

```
let on = (shape, key: { 43hr4h54j3: int }, func) -> ...
```

This has pretty good type safety, as the function will only accept keys with said field, which is hidden from the user.

This doesn't have quite as good safety guarantees as a true algebraic data type, i.e. in the form $databool = True|False$, but is certainly safer than say C enums, preventing mistakes such as using incorrect argument order.

4.7.4 Recursive Bindings

Due to the fact that everything is immutable in Vec3, the simplest way to ensure Turing completeness is to allow for recursive bindings (i.e. functions that call themselves).

This is a powerful feature as it allows for the definition of functions that operate on recursive data structures, such as trees and lists.

The type inference algorithm was modified to support recursive bindings, as the standard algorithm would not be able to infer the type of a recursive due to the fact that the binding would not be in the environment when the type of the function was inferred (all functions are lambdas, and therefore are not assigned to a binding until after declaration).

Hence, recursive functions were introduced as a separate statement in the grammar of the language, and the type inference algorithm was modified to support them by adding the binding to the environment before inferring the type of the function.

4.7.5 Constraints

Due to the restrictiveness of the standard Hindley-Milner type inference algorithm, it is not possible to support ad-hoc polymorphism (i.e. overloading) without some modification.

For example, OCaml(OCaml, 2024) does not support ad-hoc polymorphism and instead uses, for example, the $+$ operator for integer addition and the $+.$ operator for float addition, which is not ideal for this language as it would be unintuitive for a mathematician.

Examples of ML style languages that do support ad-hoc polymorphism are F#, which uses static member functions on types to achieve operator overloading(Microsoft, 2024), and Haskell, which uses type classes (Haskell, 2024) (constructs that define behaviour for a type, similarly to interfaces in object-oriented languages).

The way this issue was solved in Vec3 is by introducing the concept of *Constraint types*, which could be likened to a slightly less powerful version of type classes in Haskell.

A constraint is a type that is defined by a type variable, and a function of type *Type -¿ bool*.

During unification, if a type is unified with a constraint type, then the function is called with the type, and if it returns true then the unification is successful and the type variable that the constraint holds is unified with the type.

For example, consider the type of the $+$ operator:

```
(+) :: Constraint (a, supportsArithmetic) -> Constraint (a, supportsArithmetic) -> a
```

Then when the operator is used with say two ints, the first constrain would be unified with the type *int* (as the int type passes the *isArithmetic* function), and the type variable *a* would be replaced with *int*.

The second int would then be unified with the type *int*, and the unification would be successful.

However, if the operator was used with an rational and a float, the first constraint would be unified with the type *rational*, replacing the type variable *a* with *rational*, and the second constraint would be unified with the type rational, which does not unify with float, and so the unification would fail.

This type constrain acts as a normal type, allowing for user defined ah-hoc functions, such as:

```
let double = (x) -> x + x
```

This function can be called with any type that supports arithmetic, and the type inference algorithm will infer the type of the function as *Constraint (a, supportsArithmetic) -¿ a*.

A current limitation of this system is that the user cannot define their own constraints, and the only constraints present are those built into the language (such as operators).

This is a feature that could be added in the future, but was not necessary for the current implementation of the language.

4.7.6 Vector Length Encoding

Another key feature present in the type system is the encoding of vector lengths.

The type of a vector looks like *Vector of Type * Dims*, where *Type* is the type of the elements, and *Dims* is an integer representing the number of dimensions.

This means that a vector *[1,2,3]* is inferred to be of type *Vector of int * 3*.

This allows for the type system to catch errors such as adding two vectors of different lengths, or only allowing the *cross product* function to be called on vectors of length 3.

This is a powerful feature as it allows for the type system to catch errors that would otherwise only be caught at runtime with standard type inference.

In its current state, it also allows for slight *refinement types*(Freeman and Pfenning, 1991), which are types that are dependent on values, such as the length of a vector.

Examples of this catching an otherwise runtime error is shown in the following code:

```
let a = [1,2,3]
let b = [1,2]
let c = a + b // Error: Vectors must be of the same length

let d = [1,2,3]
let e = d[4] // Error: Index out of bounds
```

The latter example is currently very simple, and only catches out of bound errors during indexing with constant values, but could be extended to catch more complex errors in the future.

One thing to note however is that the dimensions of a vector are lost fairly easily, for example during *cons*, as it is not powerful enough to infer the length of the resulting vector.

Having full dependent types would solve this issue, but would be overkill for this language, and would make the type system much more complex (likely requiring a theorem prover and types as values).

4.7.7 Function Purity

The purity of a function is also determined during type inference, with the type of a function being inferred as pure if it is made up of only pure functions, with the base pure functions being built in.

This allows for, for example, the *plotFunc* (4.11) ensuring that only pure functions can be passed to it, preventing a user from passing a function that has side effects (which would likely cause a runtime error).

It also allows for easy dead code elimination, as a call to a function that has no side effects can be removed if the result is not used.

4.8 Optimisation

Before compilation, the AST is optimised by removing dead code and constant folding.

4.8.1 Dead code elimination

Dead code elimination is performed on the AST by removing any statements that are not used.

For example, if a variable is declared but never used, the variable declaration is removed or if an expression is calculated but never used, the expression is removed.

This is accomplished by checking every statement in the AST. If it's an expression node with no side effects, it is removed if it is not used.

If it's a binding node, the rest of the AST is checked for uses of the variable.

If the variable is not used, the binding is removed.

This process is repeated until no more dead code can be removed, allowing for long chains of dead code to be removed (for example if a variable is used in a function that is never called, the function would first be removed and then the variable).

4.8.2 Constant folding

Constant folding is performed on the AST by evaluating constant expressions at compile time, such as $2 + 2 \rightarrow 4$.

This is accomplished using the initial interpreter implementation, which recursively evaluates the AST and replaces constant expressions with their evaluated value.

Only constants are evaluated, and thus no variable resolution is performed due to the cost of this operation.

4.9 Bytecode Compiler

4.10 Virtual Machine

4.11 Prelude

A prelude is implicitly included in every program, which contains some useful functions defined in the language, as well as wrappers for the built-in functions of the Virtual Machine.

Notable functions include:

- *map*, *fold* and *filter* functions for lists.

- *range* function for generating a list of numbers.

- *sqrt*, *cubeRoot* which are specialisations of the *root* function.

- *head*, *tail* and *len* functions for lists.

- *findIntegral* function for finding the integral of a function.

We felt it was useful implementing these in-language functions as it allows for more concise and readable code, as well as showcasing the power of the language.

4.12 Plotting

The plotting system is implemented using ScottPlot(**?**), a plotting library for .NET.

The functionality is exposed to the user through 3 built-in functions: *plot*, *plotFunc* and *plotFuncs*.

plot takes in a record of configuration options of the following type:

```
type PlotConfig = {
    title: string,
    XValues: [float],
    YValues: [float],
    ptype: "bar" | "scatter" | "signal",
}
```

The resulting plot is then displayed in a separate window based on these configuration options.

TODO: IMAGES OF PLOTS

The *plotFunc* function takes in a string title and a pure function of type *float -¿ float*. The function is then plotted on the graph with an infinite range of x values.

Optionally, the user can also specify two more float values, *start* and *end*, in which case the integral of the function is calculated and displayed on the graph.

TODO: IMAGES OF PLOTS

The *plotFuncs* function takes in a string title and a list of pure functions of type *float -¿ float*.

This allows for multiple plots to be placed on the same window, which we felt was valuable for comparing functions or plotting derivatives.

4.13   Drawing

As well as plotting, the user also has the option of drawing arbitrary shapes on a canvas, and attaching event listeners to them.

This is done by means of the *draw* function, which takes in a record of configuration options of the following type:

```
type DrawConfig = {
    x: float,
    y: float,
    width: float,
    height: float,
    color: string,
    shape: "rectangle" | "circle",
    trace: bool,
}
```

Or a list of the above record type, allowing for multiple shapes to be drawn on the same canvas.

The *draw* function then returns a unique identifier for the shape, which can be used to attach event listeners, allowing for movement of the shape through key presses.

The following example attached event listeners to a shape which moves it left and right following the *cos* curve:

```
on(id, Keys.Right, (state) -> { x = state.x + 10.0, y = cos(state.x) * 10.0 + 100.0 })
on(id, Keys.Left, (state) -> { x = state.x - 10.0, y = cos(state.x) * 10.0 + 100.0 })
```

Where keys is a record defined in the prelude of the language (see sec:prelude).

Additionally, the *trace* option allows for the shape to leave a trail behind it, which can be useful for animations or visualising movement.

TODO: IMAGES OF DRAWING

4.14   Notebook View

4.15   Code architecture

Chapter 5

Discussion, conclusion and future work

Bibliography

Bierman, G., Abadi, M., and Torgersen, M. (2014). Understanding typescript. In *ECOOP 2014–Object-Oriented Programming: 28th European Conference, Uppsala, Sweden, July 28–August 1, 2014. Proceedings 28*, pages 257–281. Springer.

Freeman, T. and Pfenning, F. (1991). Refinement types for ml. In *Proceedings of the ACM SIGPLAN 1991 conference on Programming language design and implementation*, pages 268–277.

Garcia, R., Clark, A. M., and Tanter, É. (2016). Abstracting gradual typing. *ACM SIGPLAN Notices*, 51(1):429–442.

Haskell (2024). Haskell documentation. https://www.haskell.org/documentation/.

Microsoft (2024). F# documentation. https://docs.microsoft.com/en-us/dotnet/fsharp/.

Milner, R. (1978). A theory of type polymorphism in programming. *Journal of computer and system sciences*, 17(3):348–375.

Morris, J. G. and McKinna, J. (2019). Abstracting extensible data types: or, rows by any other name. *Proceedings of the ACM on Programming Languages*, 3(POPL):1–28.

OCaml (2024). Ocaml documentation. https://ocaml.org/docs/.

Sulzmann, M. F. (2000). *A general framework for Hindley/Milner type systems with constraints*. Yale University.

Timany, A., Krebbers, R., Dreyer, D., and Birkedal, L. (2024). A logical approach to type soundness. *Journal of the ACM*, 71(6):1–75.

Appendix A

Contributions

50/50

Appendix B

Testing

B.1 Lexer testing

B.2 Parser testing

B.3 Expression testing

B.4 Variable assignment testing

B.5 Function testing

B.6 GUI testing

B.7 Plot testing

Appendix C

Other stuff